

# Homework Assignment #4

November 3, 2023

## 1 Problem 1: Letter Recognition (Adapted from Bertsimas 22.3)

One of the most widespread and classical applications of machine learning is to do optical character/letter recognition. In this problem, we have four possible labels of each observation, namely whether the observation is the letter A, B, P, or R. Hence this is a **multi-class classification problem**.

The file `Letters_train.csv` and `Letters_test.csv` contains 2181 and 935 observations, each of which corresponds to a certain image of one of the four letters A, B, P and R. For each image, we have available certain **attributes of the image** in terms of pixels, as well as **which of the four letters the image is**. These features are described in Table 1. Note that **feature selection is not required for this problem**.

### 1.1 a (25 points)

To warm up, we'll start by predicting whether or not the letter is "B". To do so, first create a new variable called `isB` in your dataset, which takes value "Yes" if the letter is B, and "No" if the letter is not B.

i) Before building any models, first consider a baseline method that always predicts the most frequent outcome, which is "not B". What is the accuracy of this baseline method on the test set?

ii) Construct a logistic regression model to predict whether or not the letter is a B, using the training set to build your model. (Remember to not use the the variable letter as one of the independent variables in any of the models in this part, since it is not a valid feature.) What is the accuracy of your logistic regression model on the test set, using a threshold of  $p = 0.5$  ?

iii) What is the AUC of your logistic regression model?

iv) Construct a CART tree to predict whether or not a letter is a B, using the training set to build your model. Select the `ccp.alpha` value for the tree through cross-validation, and explain how you did the cross-validation and how you selected the `ccp.alpha` value. What is the accuracy of your CART model on the test set?

vi) Compare the accuracy of your logistic regression and CART models. Which one performs best on the test set? For this application, which do you think is more important: interpretability or accuracy?

## 1.2 b (25 points)

Now let us move on to the original problem of interest, which is to predict whether or not a letter is one of the four letters A, B, P or R. The variable in the dataset which you will try to predict is letter.

i) In a multi-class classification problem, a simple baseline model is to predict (for every observation) the most frequent class over all of the classes. For this problem, what does the baseline method predict, and what is the baseline accuracy on the test set?

ii) Now construct a CART model to predict letter using the training data. Again, select the `ccp.alpha` value for the tree through cross-validation. Again, explain how you did the cross-validation and how you selected the `ccp.alpha` value. What is the test set accuracy of your CART model?

iii) Now build a Random Forest model, and this time use cross-validation to select the `max_features` value for the Random Forests method. Explain how you did the cross validation and how you selected the `max_features` value. What is the test set accuracy for your Random Forest model?

iv) Compare the test set accuracy of your CART and Random Forest models. Which model would you recommend for this problem? For your chosen model, use bootstrapping to construct a **confidence interval** for its accuracy.

Variable	Description
<b>letter</b>	The letter that the image corresponds to (A, B, P or R).
<b>xbox</b>	The horizontal position of where the smallest box enclosing the letter shape begins.
<b>ybox</b>	The vertical position of where the smallest box enclosing the letter shape begins.
<b>width</b>	The width of this smallest box.
<b>height</b>	The height of this smallest box.
<b>onpix</b>	The total number of “on” pixels in the character image.
<b>xbar</b>	The mean horizontal position of all of the “on” pixels.
<b>ybar</b>	The mean vertical position of all of the “on” pixels.
<b>x2bar</b>	The mean squared horizontal position of all of the “on” pixels in the image.
<b>y2bar</b>	The mean squared vertical position of all of the “on” pixels in the image.
<b>xybar</b>	The mean of the product of the horizontal and vertical position of all of the “on” pixels in the image.
<b>x2ybar</b>	The mean of the product of the squared horizontal position and the vertical position of all of the “on” pixels.
<b>xy2bar</b>	The mean of the product of the horizontal position and the squared vertical position of all of the “on” pixels.
<b>xedge</b>	The mean number of edges (the number of times an “off” pixel is followed by an “on” pixel, or the image boundary is hit) as the image is scanned from left to right, along the whole vertical length of the image.
<b>xedgeycor</b>	The mean of the product of the number of horizontal edges at each vertical position and the vertical position.
<b>yedge</b>	The mean number of edges as the images is scanned from top to bottom, along the whole horizontal length of the image.
<b>yedgexcor</b>	The mean of the product of the number of vertical edges at each horizontal position and the horizontal position.

Table 1: Variables in the dataset Letters.